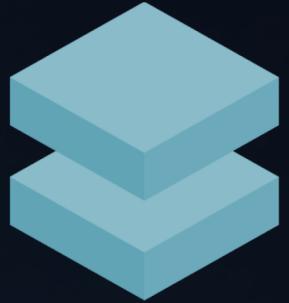




**ONE WAY**  
SOLUTION



# One Way Solution Near Real-Time ETL

Data Engineering – [Day 3]

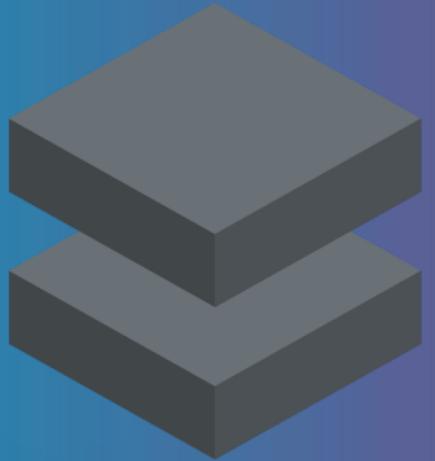


**LUAN MORENO**  
CEO & CDO  
Data Engineer & Data Platform MVP  
Confluent Certified Developer for Apache Kafka [CCDAK]

# Batch-ETL



Test Data Engineering – [Day 2]



# Data Engineering – Batch-ETL



2

Data Lake  
Spark Lifecycle  
Storage Solutions for Spark  
Data Lakehouse  
Delta Architecture





# Agenda

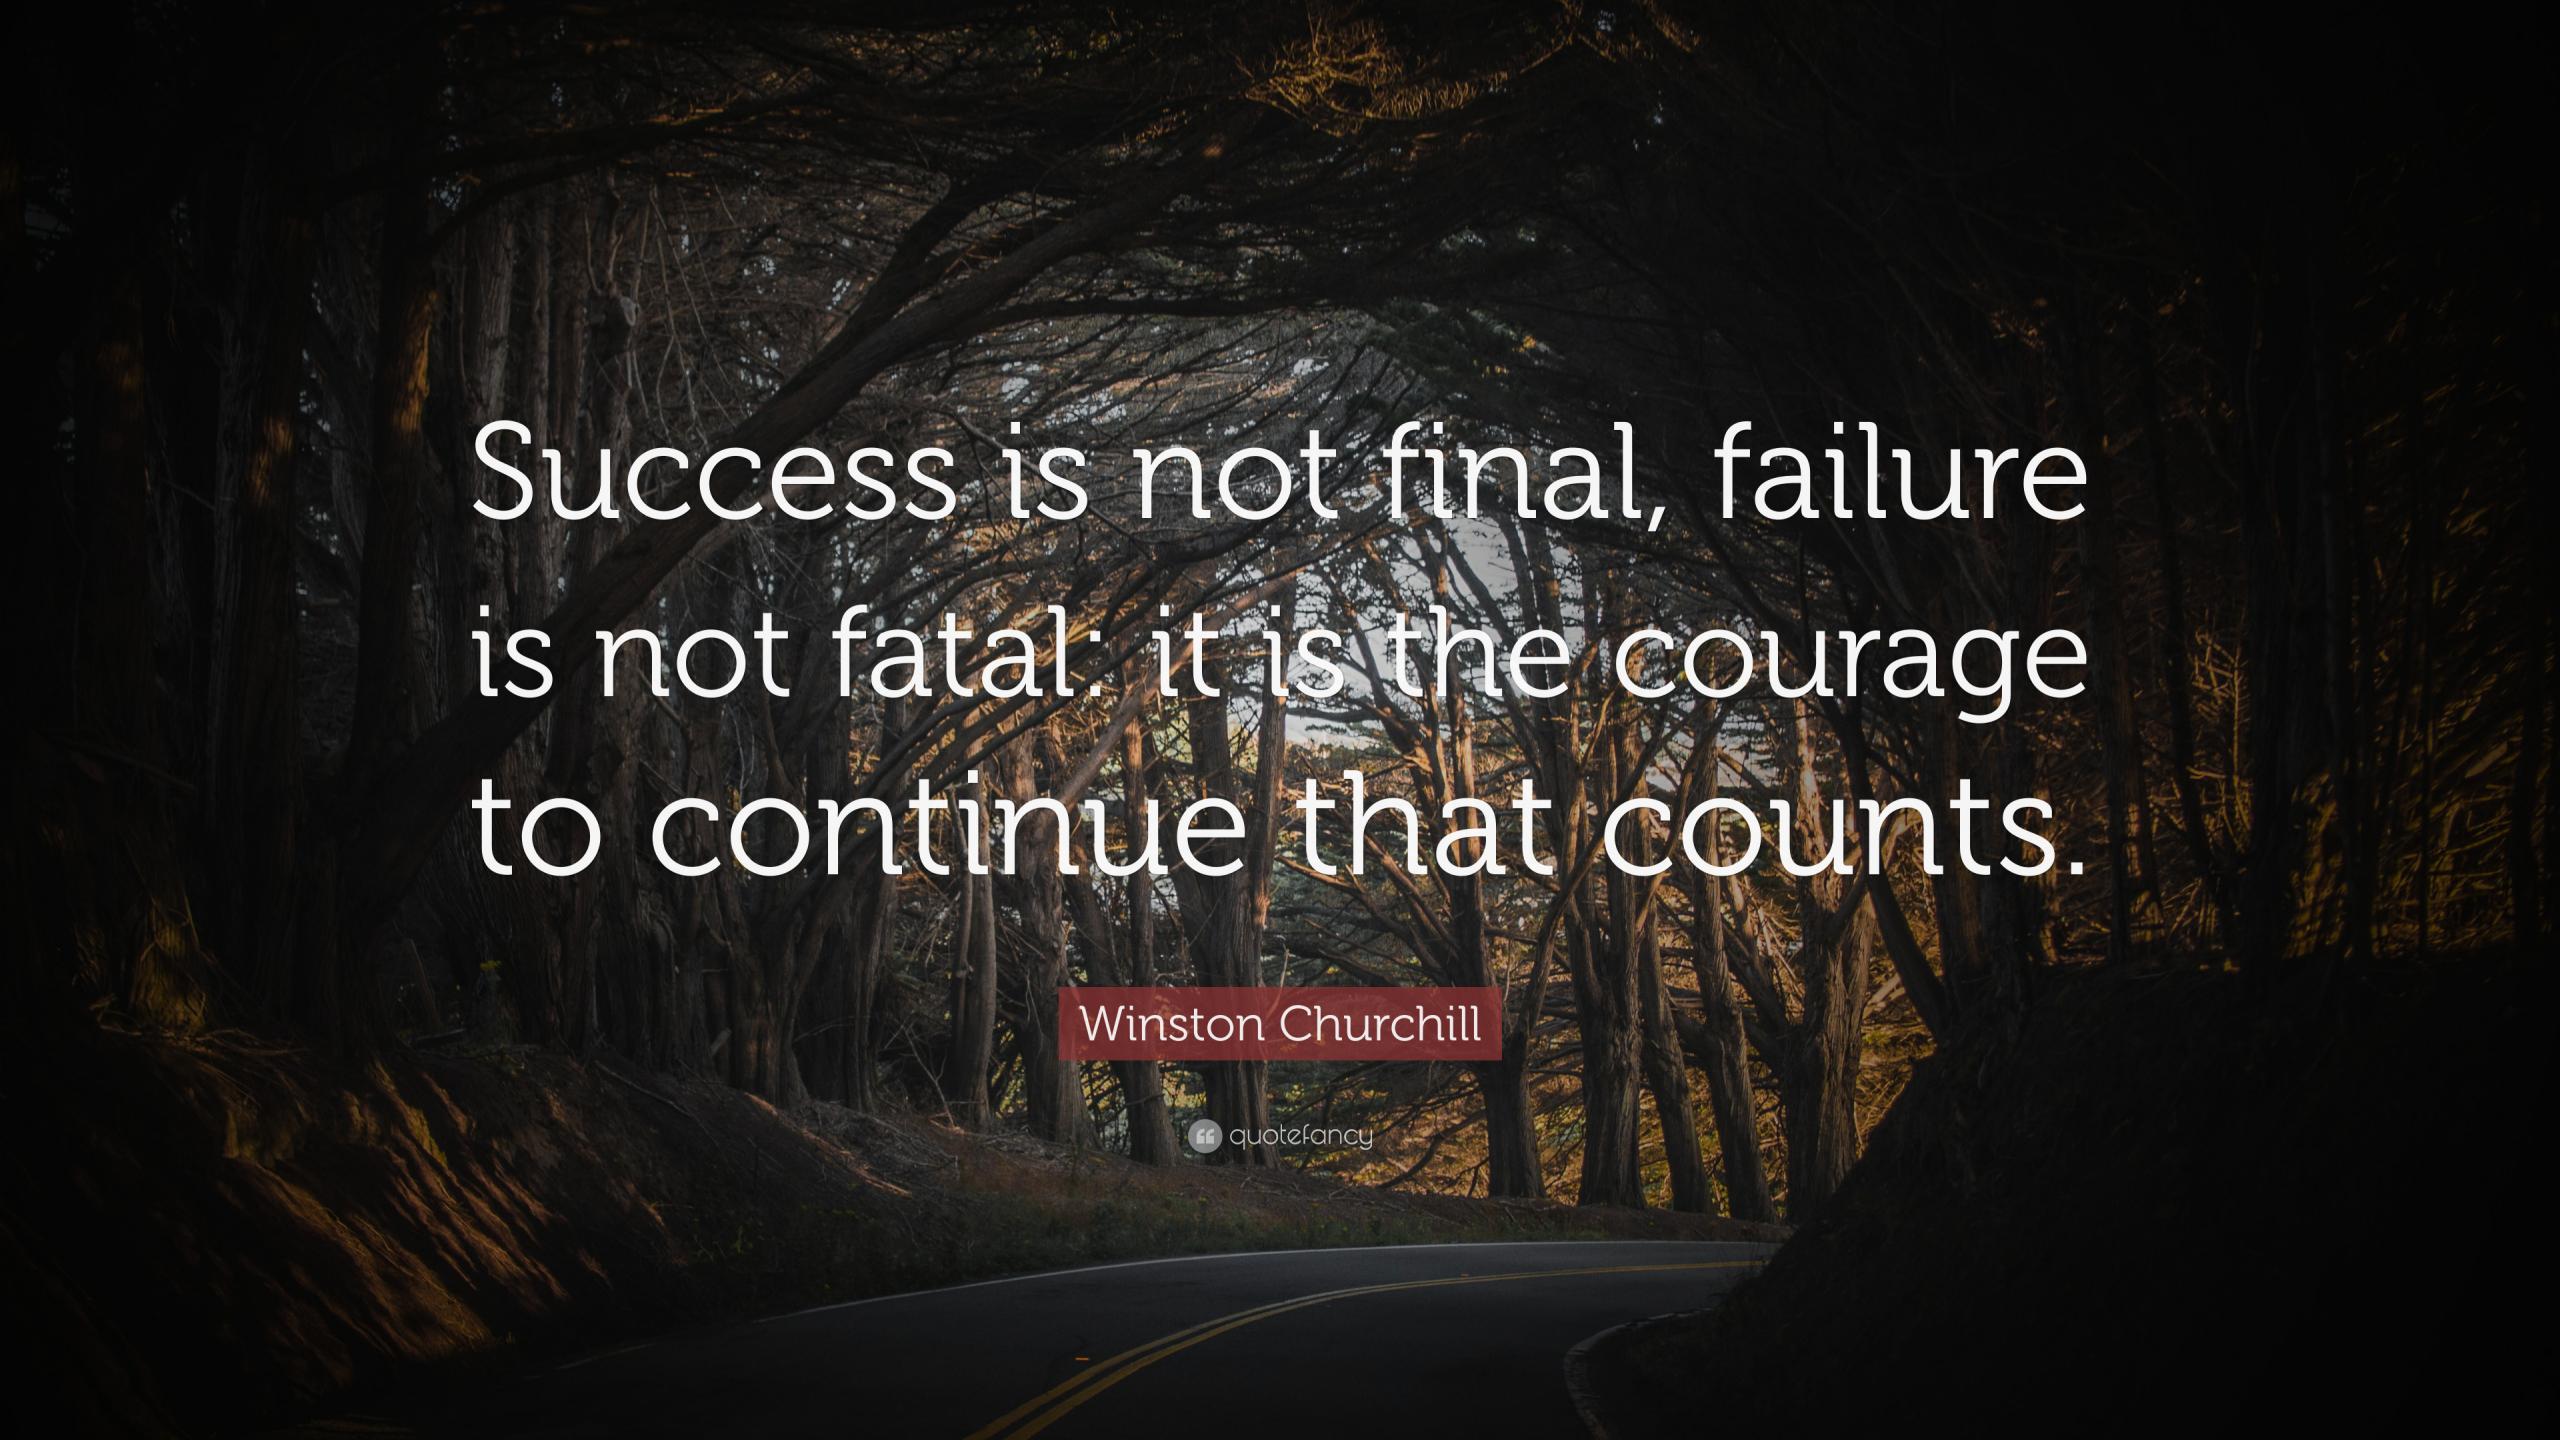
3

Event Stream Processing  
Real-Time Stream Processing Engines  
Python & Streaming Engines  
StreamingSQL  
Apache Kafka  
Spark Streaming & Structured Streaming  
Use-Cases



*One Way Solution*





Success is not final, failure  
is not fatal: it is the courage  
to continue that counts.

Winston Churchill



# Event Stream [ES]

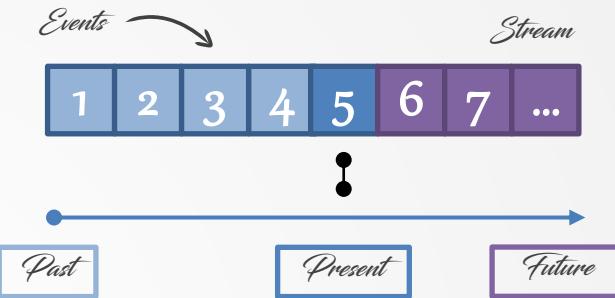


Event Stream [ES] = Representation of an Unbounded DataSet  
Unbounded = Infinite & Ever Growing



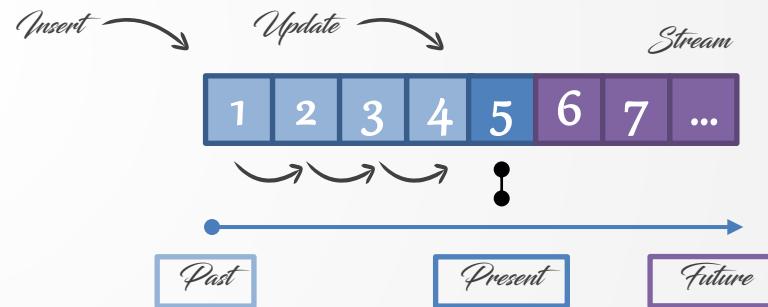
## Event Streams [Ordered]

There is an inherent notion of which events occur before or after other events. This is clearest when looking at financial events. A sequence in which I first put money in my account and later spend the money is very different from a sequence at which I first spend the money and later cover my debt by depositing money back. The latter will incur overdraft charges while the former will not. Note that this is one of the differences between an event stream and a [database table](#) [records in a table are always considered unordered](#) and the order by clause of SQL is not part of the relational model.



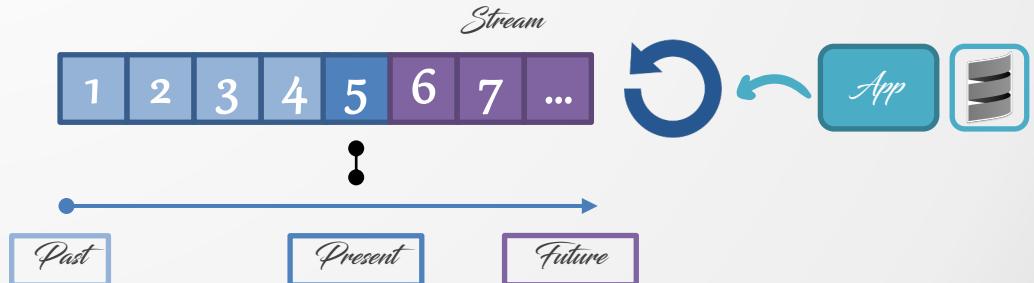
## Immutable Data Records

Events, once occurred, can never be modified. A financial transaction that is cancelled does not disappear. Instead, an additional event is written to the stream, recording a cancellation of previous transaction. When a customer returns merchandise to a shop, we don't delete the fact that the merchandise was sold to him earlier, rather we record the return as an additional event. This is another difference between a data stream and a database table. We can delete or update records in a table, but those are all additional transactions that occur in the database, and as such can be recorded in a stream of events that records all transactions. If you are familiar with binlogs, WALS, or redo logs in databases you can see that if we insert a record into a table and later delete it, the table will no longer contain the record, but the redo log will contain two transactions - the insert and the delete.



## Event Streams [Replayable]

This is a desirable property. While it is easy to imagine nonreplayable streams (TCP packets streaming through a socket are generally nonreplayable), for most business applications, it is critical to be able to replay a raw stream of events that occurred months (and sometimes years) earlier. This is required in order to correct errors, try new methods of analysis, or perform audits. This is the reason we believe [Kafka made stream processing so successful in modern businesses](#)—it allows capturing and replaying a stream of events. Without this capability, stream processing would not be more than a lab toy for data scientists.

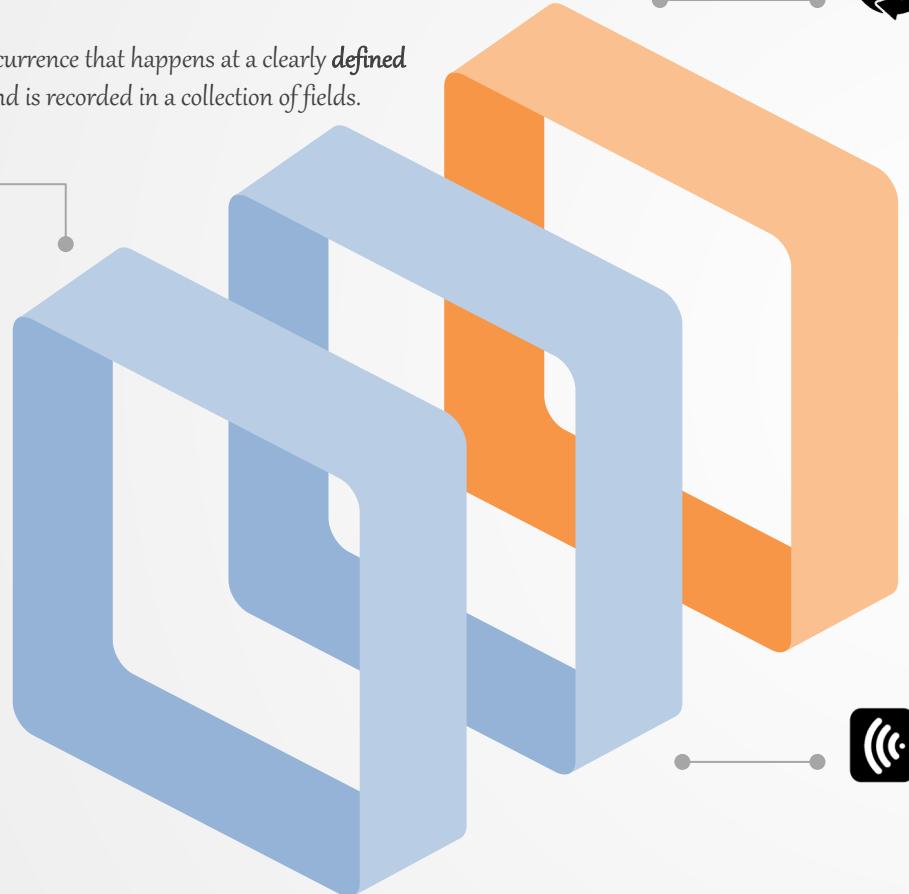


# Event Stream Processing [ESP]



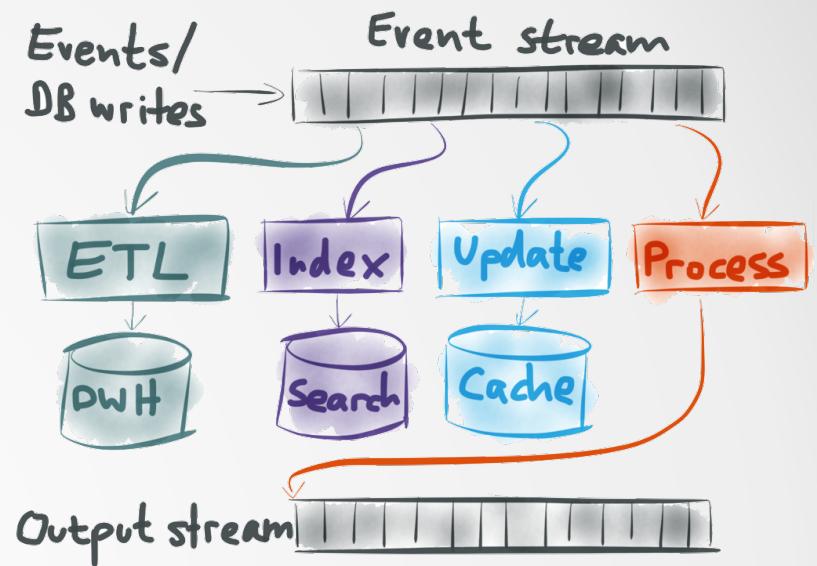
## Event

any occurrence that happens at a clearly **defined time** and is recorded in a collection of fields.



## Processing

the act of **analyzing data** and perform data analytics on top of that.



## ESP for Processing Changes

- Store Data Reliability
- Process Incoming Data
- Perform Queries & Analytics
- Push Results Immediately to Subscribers

# Real-Time Stream Processing [Engines]



## Open-Source Platform [OSS]

- Apache Kafka
- Apache Spark 
- Apache Apex
- Apache Flink
- Apache Storm
- Apache Beam



## Microsoft Azure

- HDInsight
- Azure Synapse Analytics
- Azure Stream Analytics
- Azure Functions



## Google Cloud Platform [GCP]

- Cloud DataProc
- Cloud DataFlow
- Cloud Functions



## Amazon Web Services [AWS]

- Amazon EMR
- Amazon Kinesis Data Streams
- Amazon Kinesis Data Analytics
- AWS Lambda

# Python & Streaming Engines

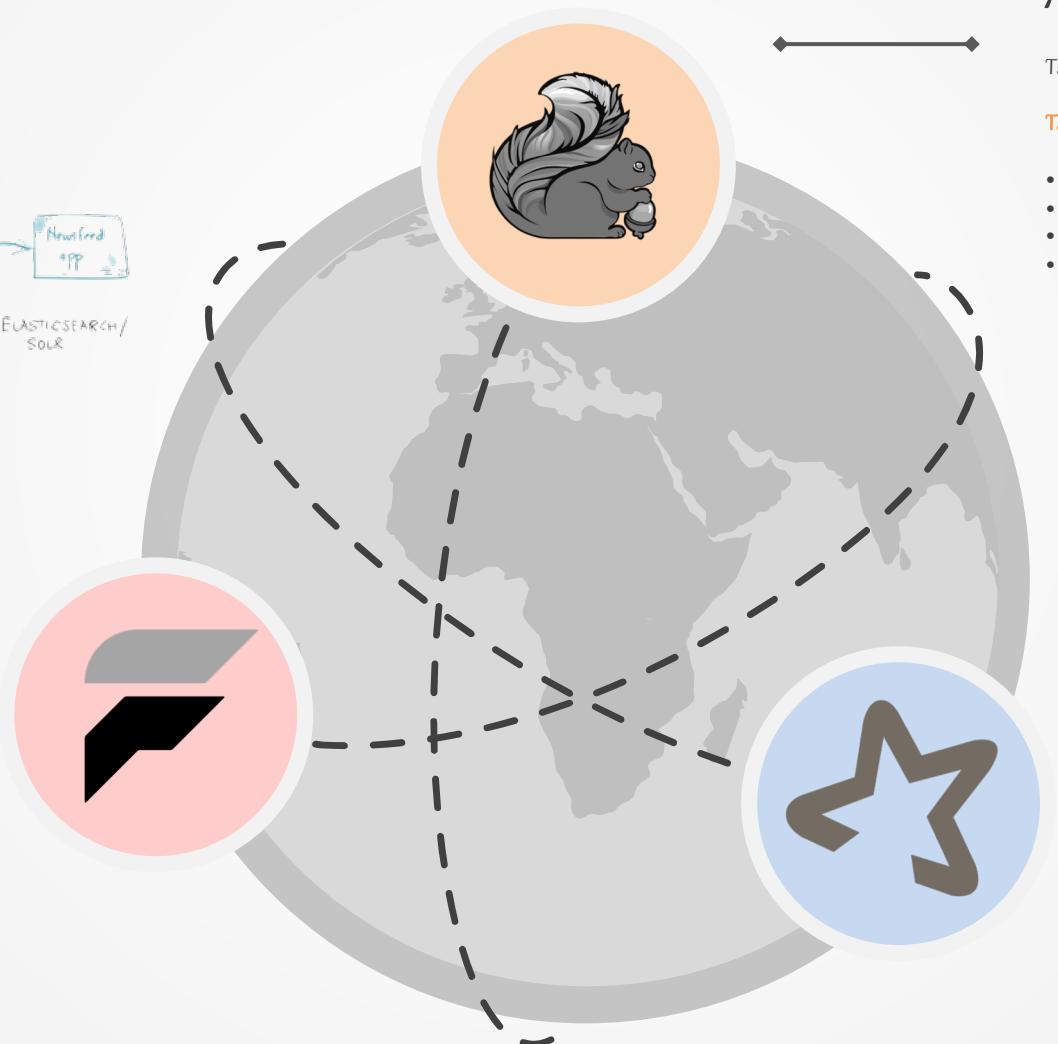


## Faust

Python Stream Processing Library

### Faust API

- Stream Processing ~ Kafka Streams, Spark, Storm, Samza, Flink
- Numpy, PyTorch, Pandas, NLTK, Django
- Tables ~ Distributed Key & Value
- Use Regular Python Dictionaries



## Apache Flink

Table API for a Unified Stream & Batch Processing Experience

### Table API

- Language-Integrated Query API for Scala & Java
- Batch & Streaming Input without Code Changes
- Fully Integrated with DataStream and DataSet APIs
- Integrated with Complex Event Processing API

## Apache Spark

PySpark ~ Python API for Apache Spark Engine

### PySpark API

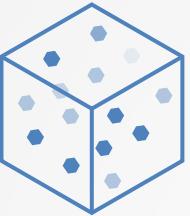
- PySpark ~ Spark DataFrame
- Distributed Table ~ Apache Spark Cluster
- SQL, Streaming, ML ~ Packages
- RDD, DStream & DataFrame

# StreamingSQL [Engines]



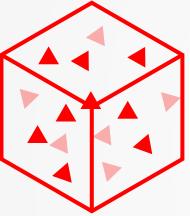
Apache Flink

Flink SQL  
2016



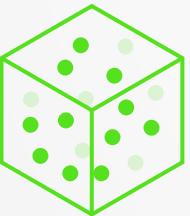
Apache Spark

Structured Streaming  
2016



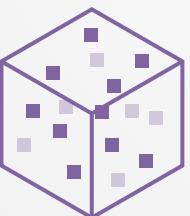
Apache Kafka

KSQL  
2017

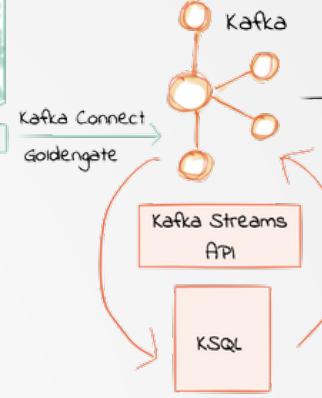
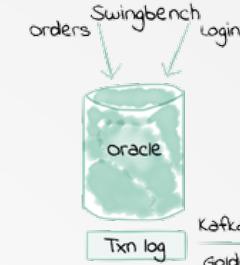


Apache Beam

Beam SQL  
2017



Streaming SQL



Data Processing

- Manipulate Streams
- Query over Continuous Flow of Data
- Select | Join | Union & Merge | Window & Aggregation
- Real-Time Analytics
- Predictions & ML

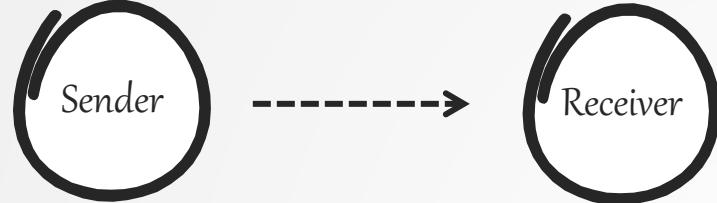


Those who dare to fail  
miserably can achieve greatly.

John F. Kennedy

# Message Delivery Guarantees [Basics]

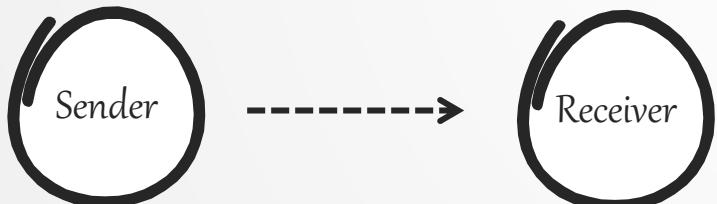
**1**



At Most Once

message may be lost

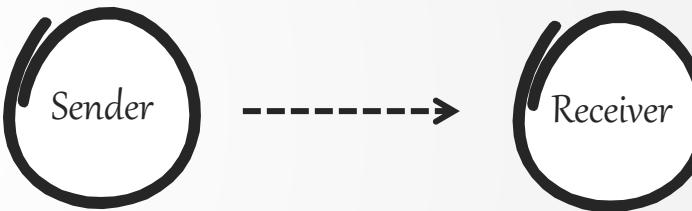
**2**



At Least Once

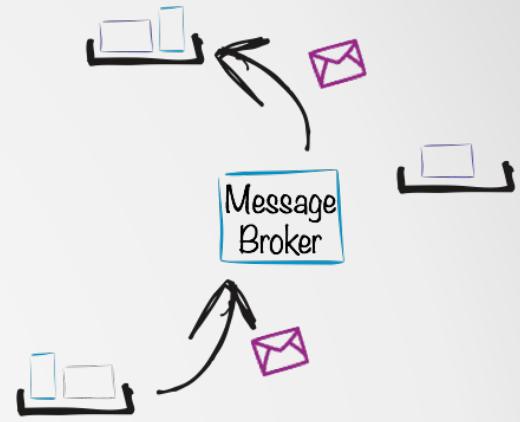
message may be duplicated  
message ordering issues

**3**



Exactly Once

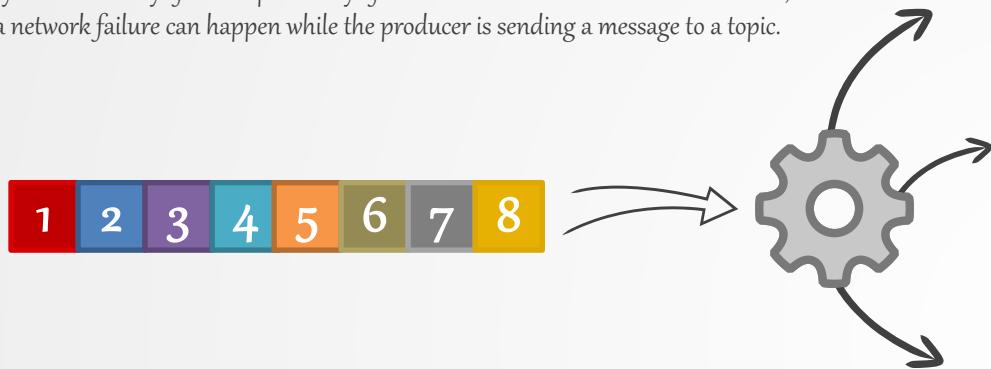
message ordering issues expensive



# Message Delivery Guarantees [Idempotent Producers & EOS]

## Messaging System Semantics

in a distributed publish-subscribe messaging system, the computers that make up the system can always fail independently of one another. Individual broker can crash, or a network failure can happen while the producer is sending a message to a topic.



### At [Most] Once

if producer does not retry when an ack times out or returns an error, the message might end up not being written to a kafka topic, and hence not delivered to consumer.



### At [Least] Once

producer receives an (ack) from broker with acks=ALL, meaning it was written once to the kafka. however, if an error arise, it might retry sending the same message, assuming that the message was not written.



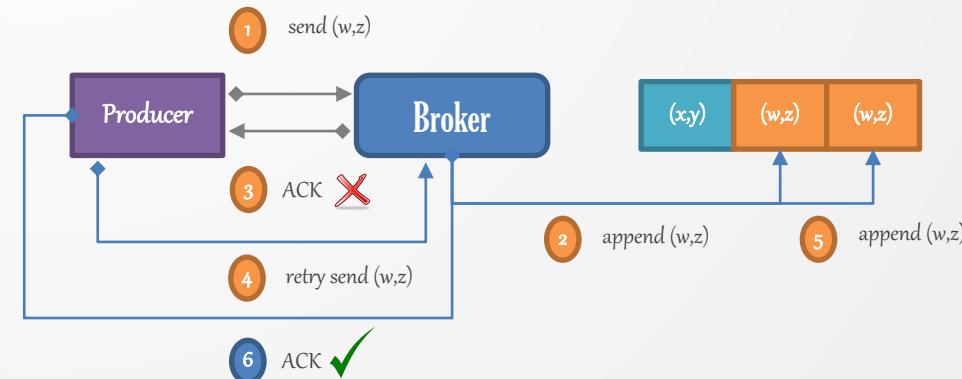
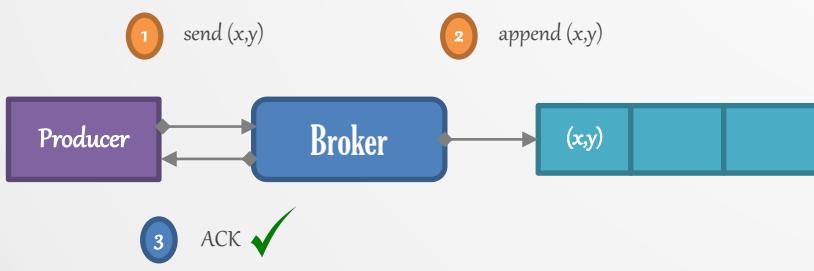
### [Exactly] Once Semantics [EOS]

even if a producer retries sending a message, it leads to the message being delivered exactly once to the end consumer. requires cooperation between the systems.



## Idempotent Producers

can be performed many times without causing a different effect than only being performed once. if retries occurs, same message won't be written again. each message will contain a sequence number which broker will use to dedup any duplicate send.



A wide-angle landscape photograph of a mountain range. The foreground is filled with dark, silhouetted pine trees. In the middle ground, a valley opens up, showing a small town with numerous houses and some snow-covered fields. The background is dominated by majestic, snow-capped mountain peaks. The sky is a clear, pale blue, suggesting either sunrise or sunset. The overall atmosphere is serene and inspiring.

Freedom lies in  
being bold.

Robert Frost

# Apache Kafka



## Apache Kafka

Open-Source Software – [Scala & Java]

7,298 Companies using Apache Kafka

+ 142 Connectors on Confluent Hub

De-Facto Streaming Platform of Fortune 100 [70%]

High-Throughput & Low-Latency Real-Time Data Feeds

Storage & Processing System using a Distributed Transaction Log



## History

Franz Kafka ~ Apache Kafka

Developed by LinkedIn

Open-Sourced in 2011

Jun Rao | Jay Kreps | Neha Narkhede ~ Confluent

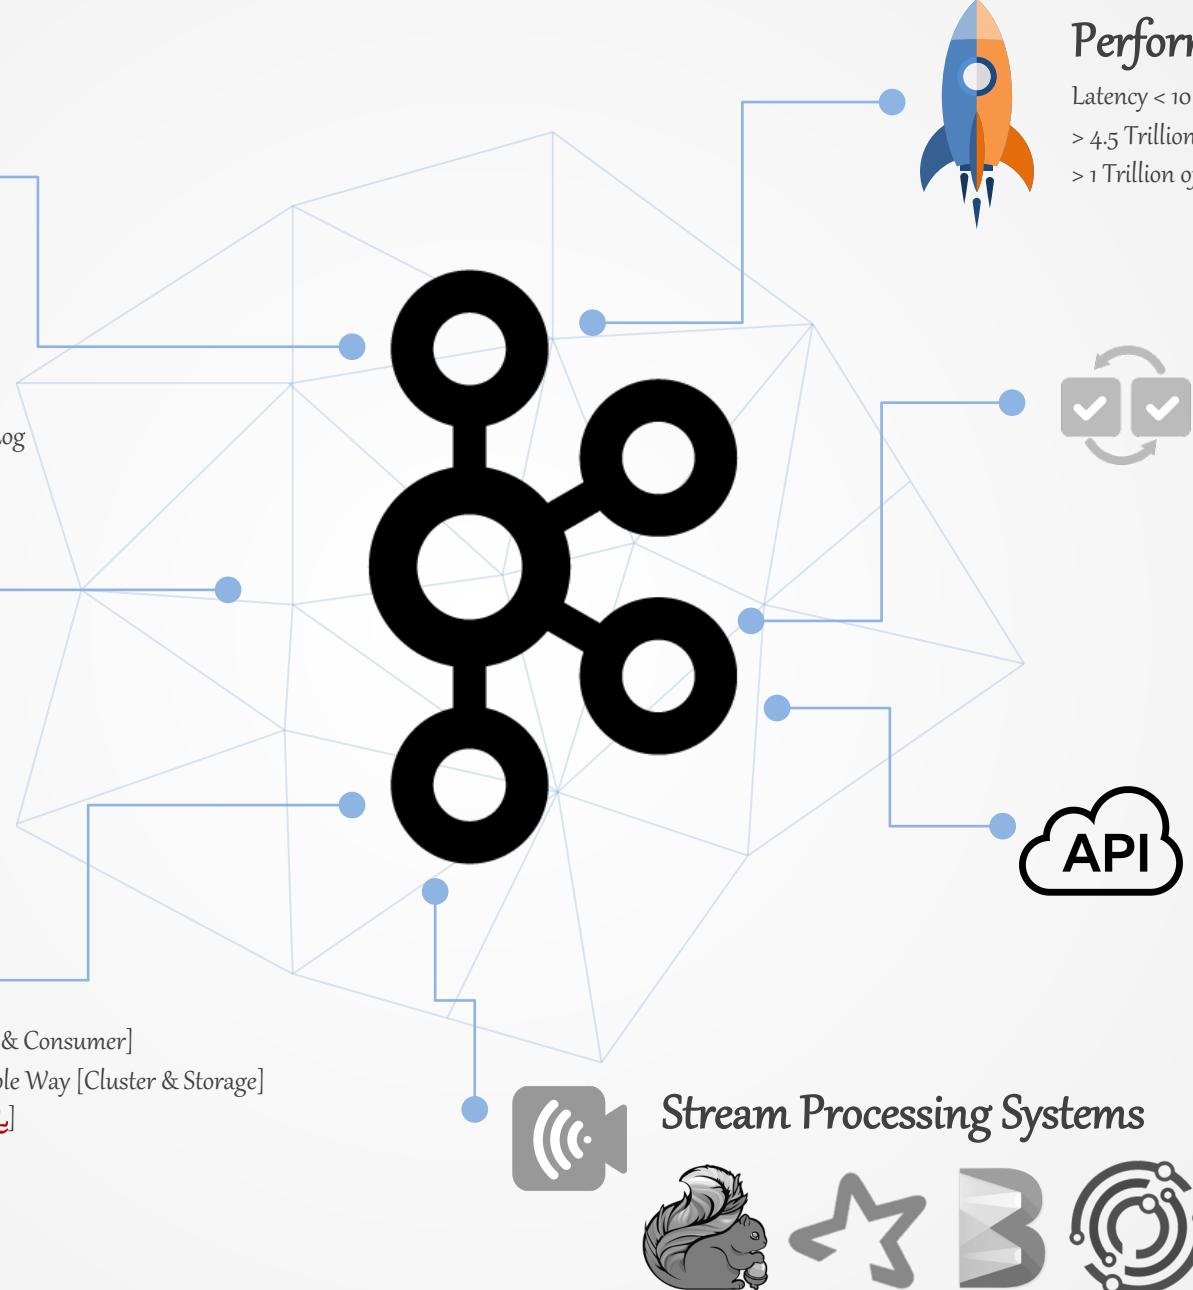


## Key Capabilities

Publish & Subscribe Streams of Records [Producer & Consumer]

Store Streams of Records in a Fault-Tolerant Durable Way [Cluster & Storage]

Process Streams of Records [Kafka Streams & KSQL]

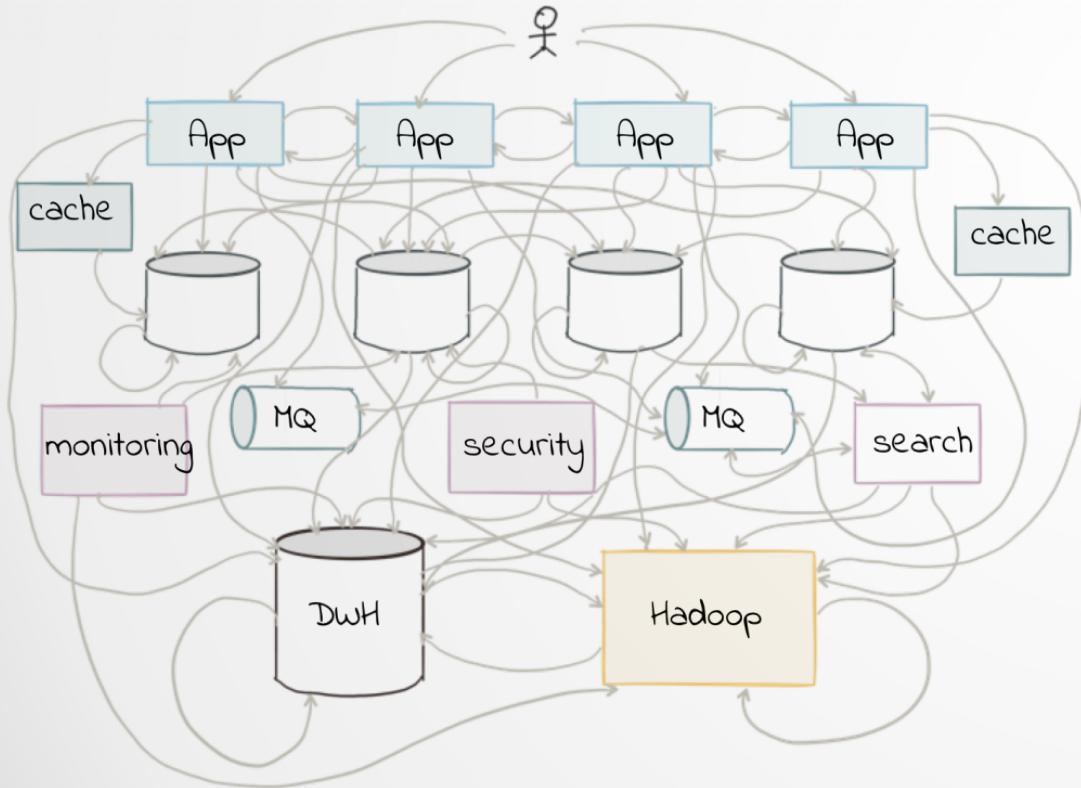


# Apache Kafka [De-Facto Streaming Platform]



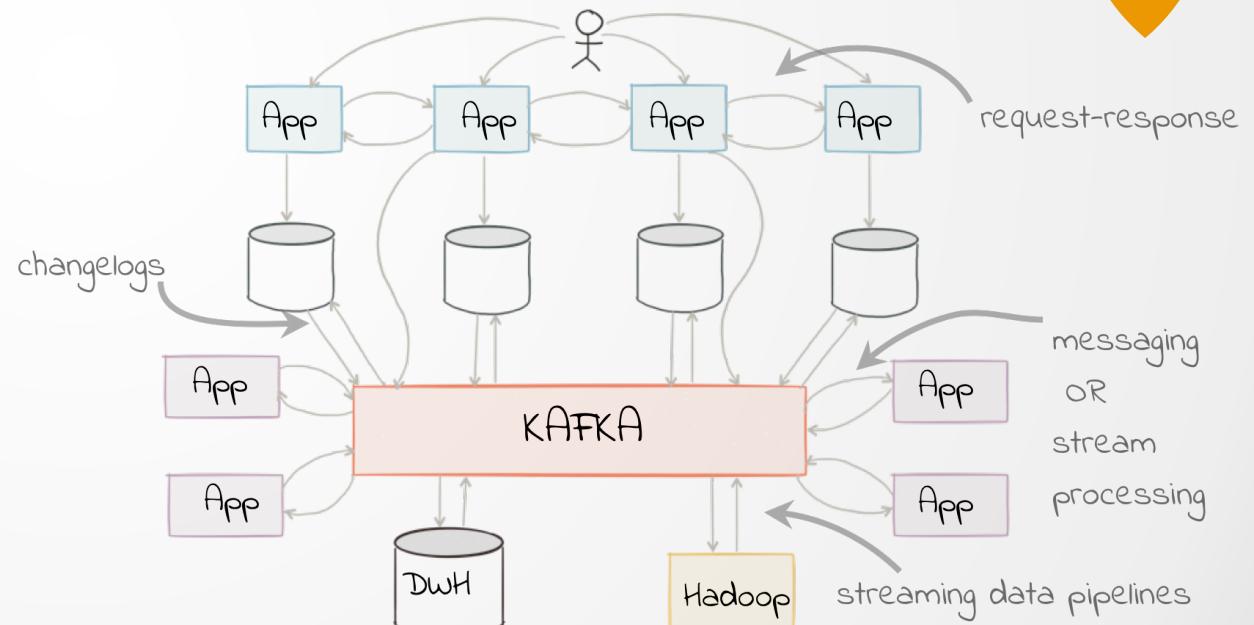
## < Apache Kafka

Data Stored in a Variety of Places  
Difficult for Data Integration  
Dramatic Business Impact  
Siloed Data



## > Apache Kafka

Centralized Data Exchange Hub  
Data Sent in Binary  
Fast & Reliable Architecture  
Commit-Log Structured  
Easy for Data Integration

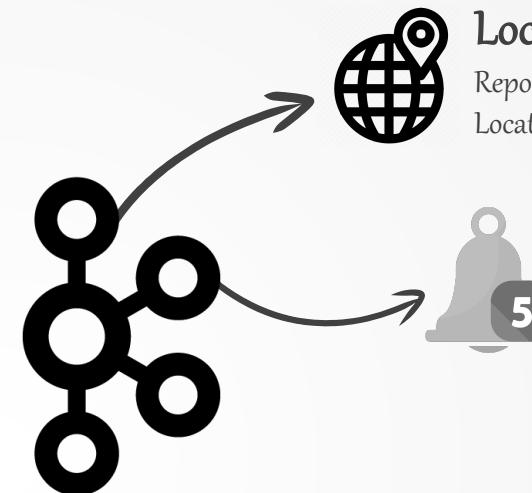


# Apache Kafka [Use-Case]



**GPS**

Global Positioning System  
Truck Location Data  
Every 20 Seconds



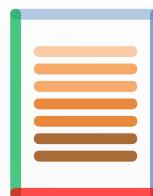
**Location Dashboard**

Report & Dashboards  
Location of Trucks



**Notification Service**

Report & Dashboards  
Location of Trucks  
Weather Conditions & Fuel Status

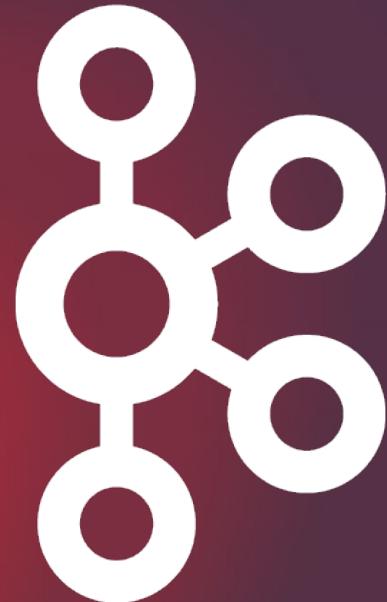


**Topic**

Topic Name – `gps_location_trucks`  
Data – `truck_id | latitude | longitude`  
Partitions – 5 [Ordered]  
Key – `truck_id`



# Getting Started with Apache Kafka



Live out of your imagination,  
not your history.

Stephen R. Covey

# Spark Streaming



## Continuous Streams of Data

High-Level Abstraction – Discretized Streams or DStream [RDD]



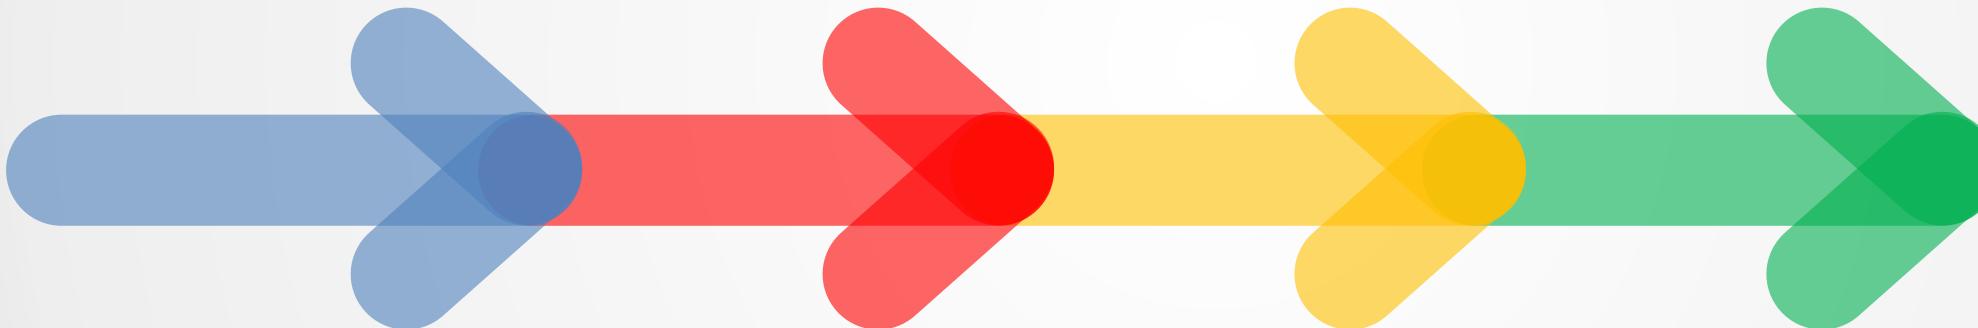
### DStreams

Continuous Series of RDDs [Immutable]  
RDD – Distributed Table In-Memory  
Store RDD in Different Partitions  
Transformation & Action – Lazy Evaluation



### DStreams Transformations

- Map | FlatMap | Filter | Repartition
- Union | Count | Reduce | CountByValue
- ReduceByKey | Join | CoGroup
- Transform | UpdateStateByKey



### Data Sources

- File Systems
- Socket Connections
- Apache Kafka
- Apache Flume
- Amazon Kinesis



### Python API

- Spark Streaming 2.4.4  
PySpark Streaming Module
- Apache Kafka [0.8.2.1]
  - Apache Flume [1.6.0]
  - Amazon Kinesis [1.2.1]

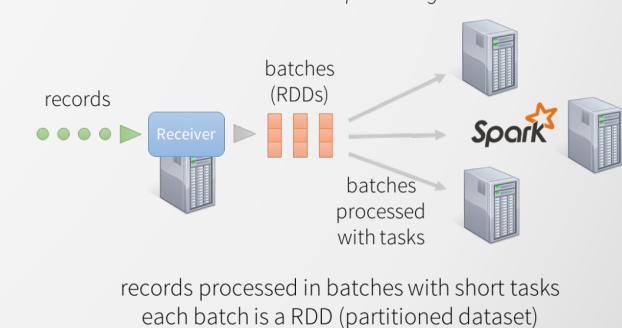


### Output Operations

- Pushed Out to External Systems  
Trigger DStream Transformation  
Java & Scala – ALL Outputs  
Python – Print | Txt Files | ForeachRDD

## Spark Streaming

discretized stream processing

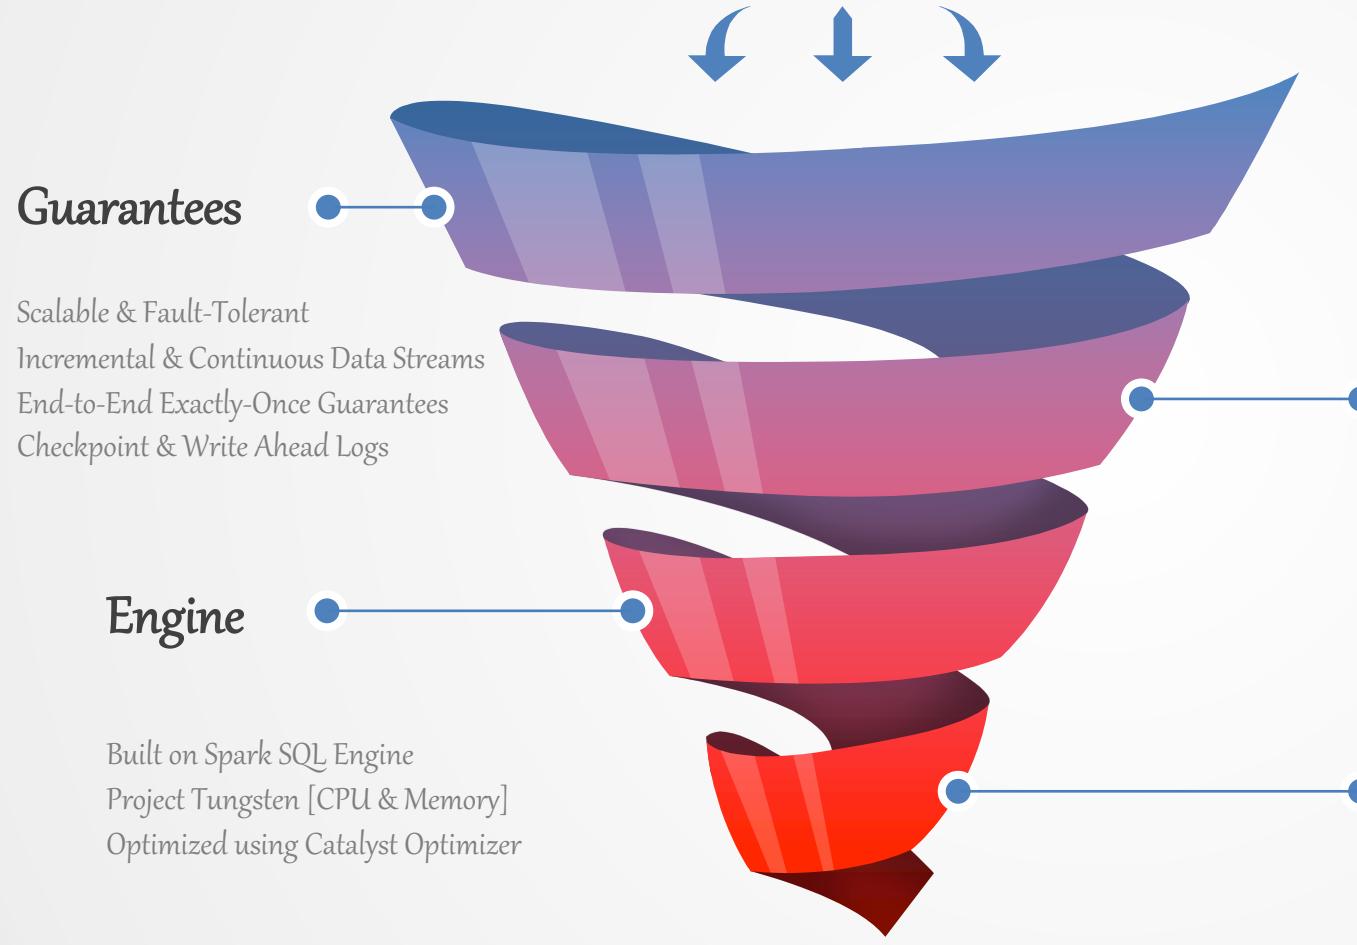


A wide-angle landscape photograph of a mountain range. The mountains in the background are heavily covered in snow, with dark, rocky peaks visible through the white. In the foreground, a valley is visible with a small, snow-dusted town at the base. The sky is a clear, pale blue, suggesting either sunrise or sunset. The overall atmosphere is serene and majestic.

Do not fear mistakes. You  
will know failure.  
Continue to reach out.

Benjamin Franklin

# Structured Streaming



## Philosophy

Treat Data Streams = **Unbounded Tables**  
Incremental Query over Streams

## Seamless

Express Streaming Computation = Batch Computation

## Processing

Uses Micro-Batch Processing Engine  
100 milliseconds with Exactly-Once Latency

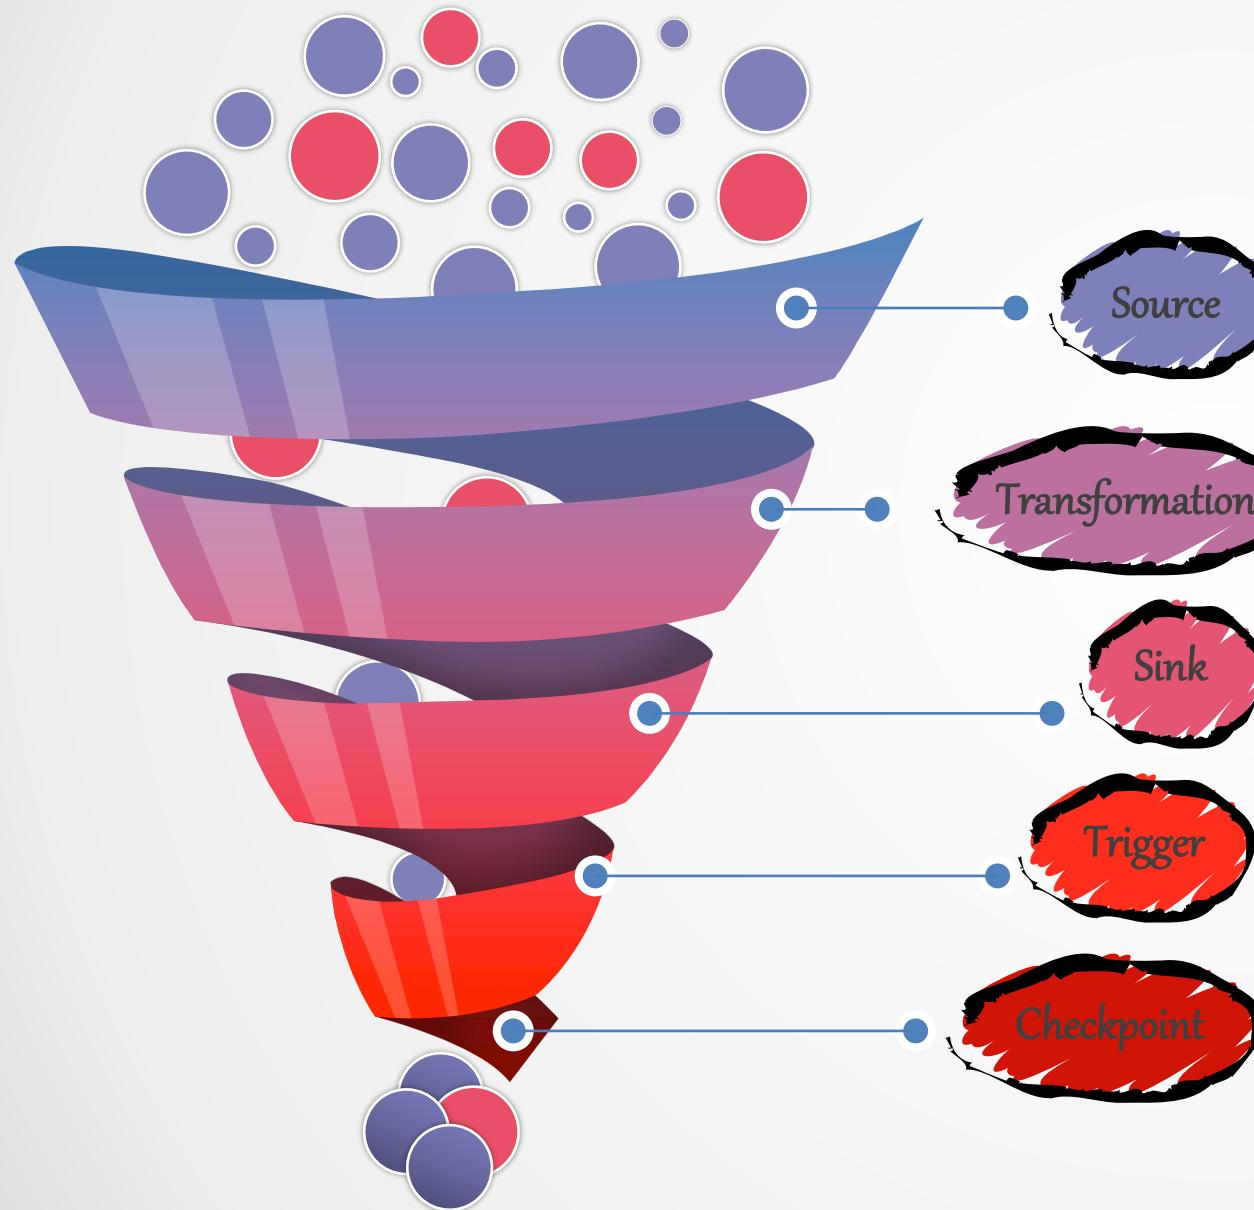
## Guarantees

Scalable & Fault-Tolerant  
Incremental & Continuous Data Streams  
End-to-End Exactly-Once Guarantees  
Checkpoint & Write Ahead Logs

## Engine

Built on Spark SQL Engine  
Project Tungsten [CPU & Memory]  
Optimized using Catalyst Optimizer

# Anatomy of a Streaming Query

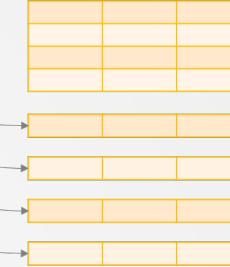


data stream as an unbounded table

Data Streams



Unbounded Table



new data in the data stream

=

new rows appended to an unbounded table

```
spark.readStream  
.format("kafka")  
.option("subscribe", "input")  
.load()
```

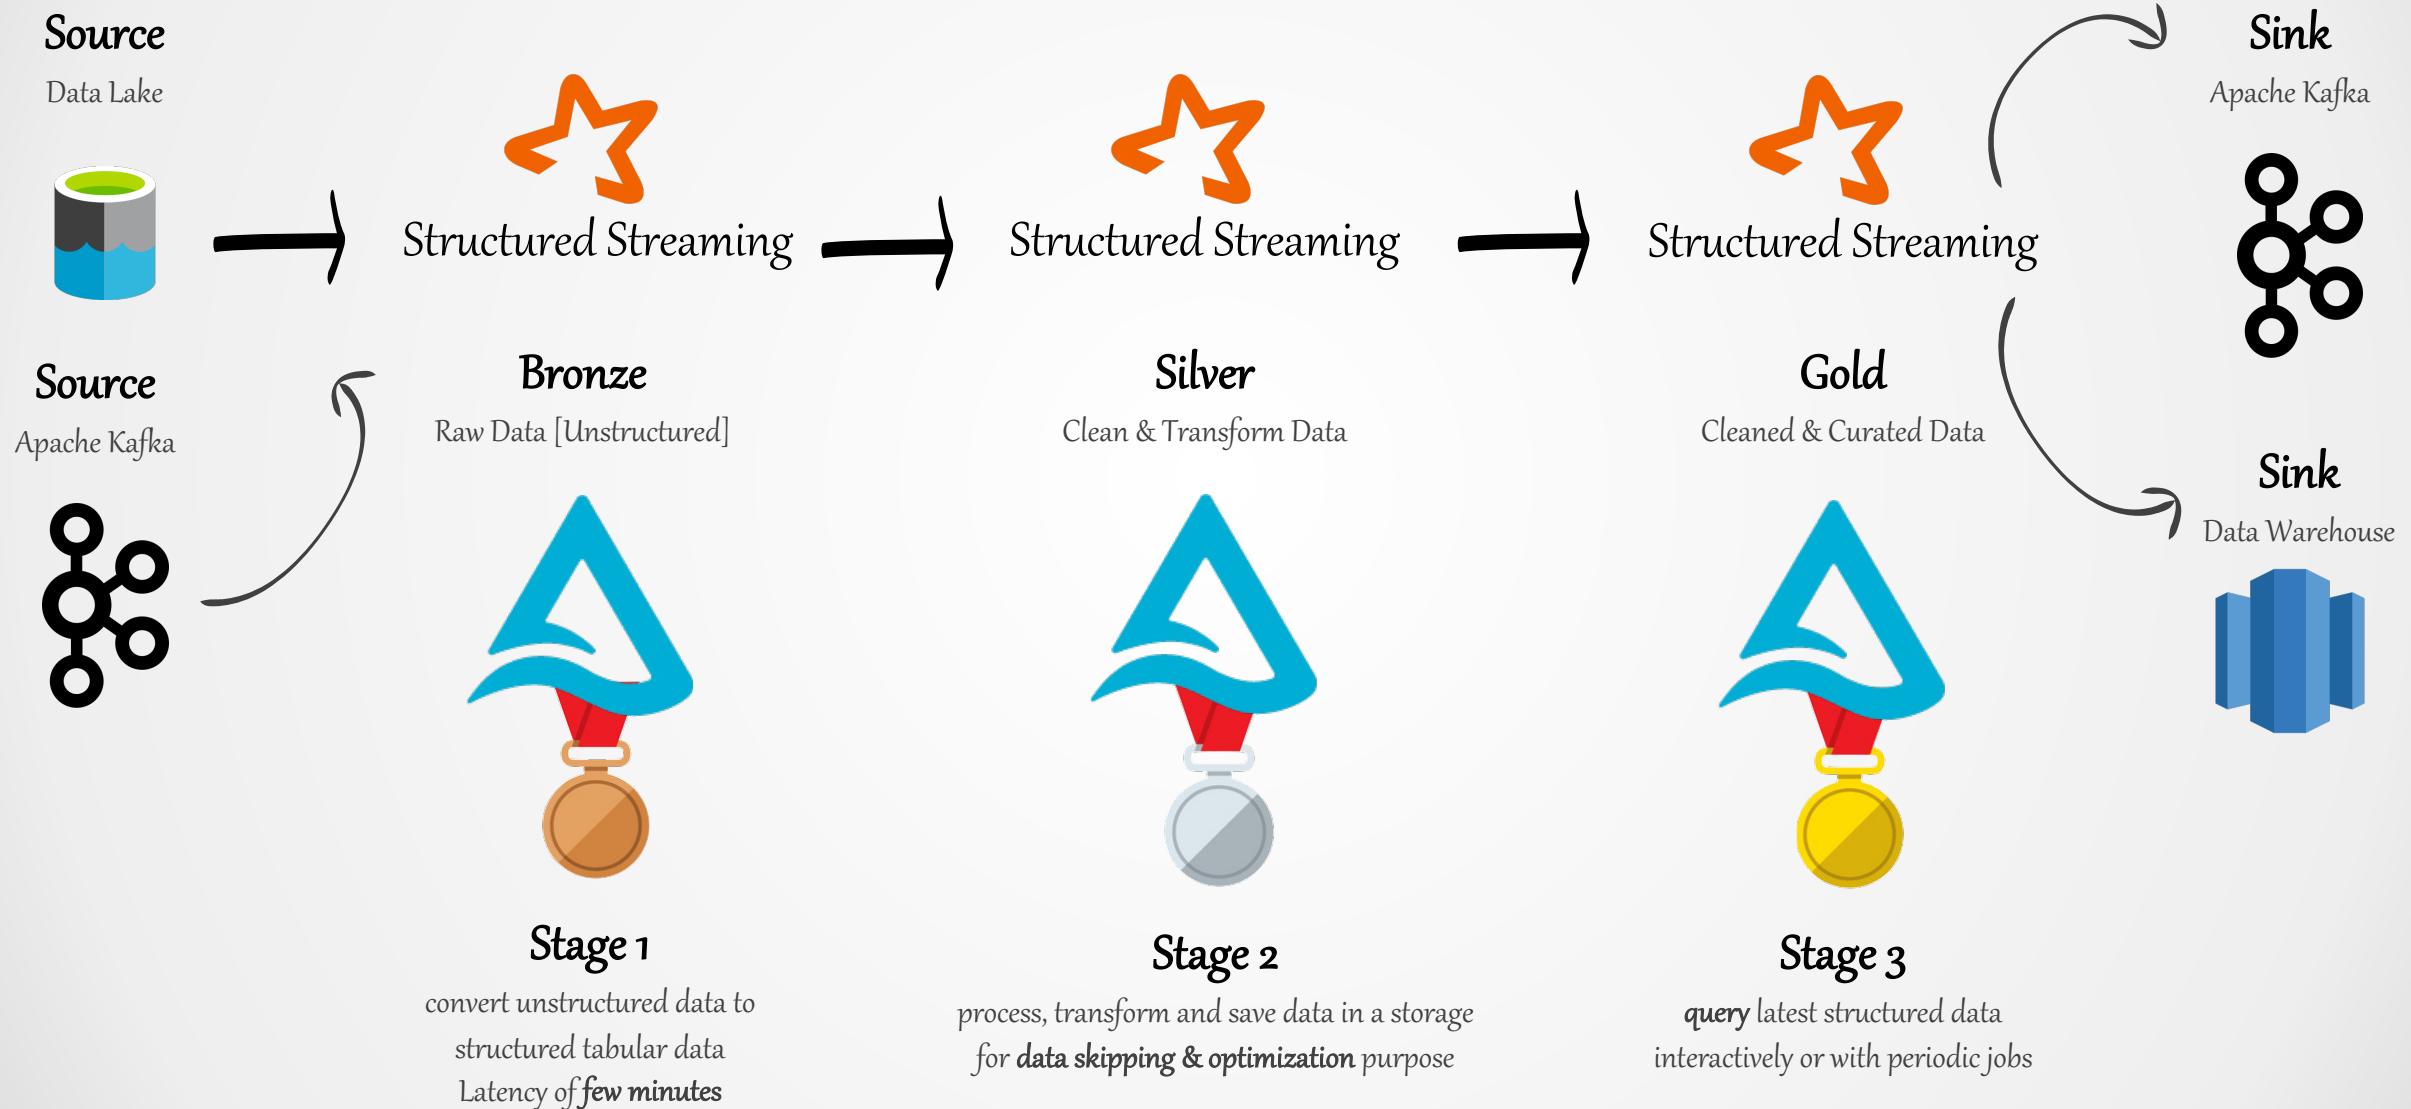
```
.groupBy('value.cast("string") as 'key)  
.agg(count("*") as 'value')
```

```
.writeStream  
.format("kafka")  
.option("topic", "output")
```

```
.trigger("1 minute")  
.outputMode("update")
```

```
.option("checkpointLocation", "...")  
.start()
```

# Structured Streaming [Use-Case] ~ Near Real-Time ETL



# Building a Data Lakehouse using Near Real-Time ETL [Data Lake & Apache Kafka]





A goal without a  
plan is just a wish.

Antoine de Saint-Exupéry

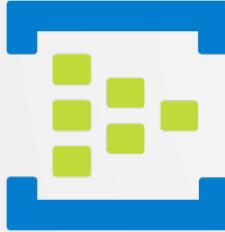
# Use-Case: Microsoft Azure



## Azure Event Hubs

Simple, Secure, & Scalable Real-Time Data Ingestion

- AMQP
- HTTPS
- Apache Kafka



## Azure Stream Analytics

Serverless Real-Time Analytics [SQL Interface]  
Machine Learning



## Azure Synapse Analytics

Fast, Flexible, & Secure Cloud Data Warehouse for Enterprises  
SQL & PolyBase Features with Fast Loading Operations



# Use-Case: Google Cloud Platform [GCP]



## Google Pub/Sub

Global Messaging & Event Ingestion  
Scale without Provisioning, Partitioning, or Load Isolation  
Expand Pipelines to New Regions Simply with Global Topics



## Cloud DataFlow

Simplified Stream & Batch Data Processing  
Apache Beam [Java | Python | SQL]

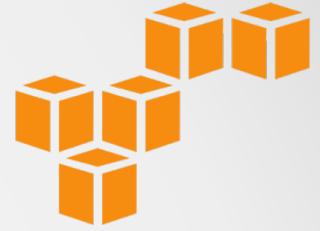


## Google BigQuery

ServerLess [SaaS], Highly-Scalable, & Cost-Effective Cloud DW  
In-Memory BI Engine & ML  
Gartner 2019 – Magic Quadrant for Data Management Solutions



# Use-Case: Amazon AWS



## Amazon Kinesis

Easily Collect, Process & Analyze Streams in Real-Time

- Kinesis Data Streams
- Kinesis Data Firehose



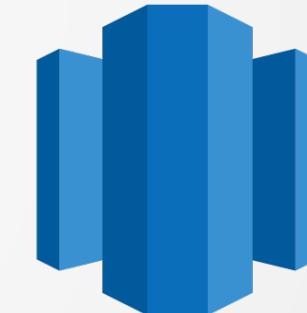
## Amazon Kinesis Data Analytics

Analyze Streaming Data  
Query Streams of Data



## Amazon Redshift

Fast, Simple, Cost-Effective Modern Data Warehouse  
MPP | ML | Result Caching & S3 Query Access

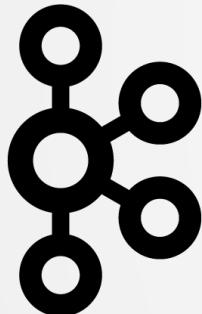


# Use-Case: Open-Source Software [OSS]



## Apache Kafka

Real-Time Data Pipelines & Streaming Apps  
Horizontally Scalable, Fault-Tolerant & Wicked Fast



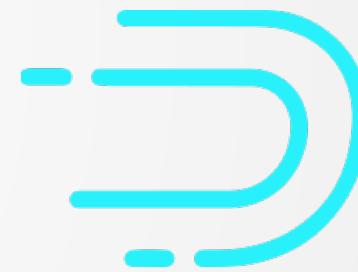
## Apache Spark

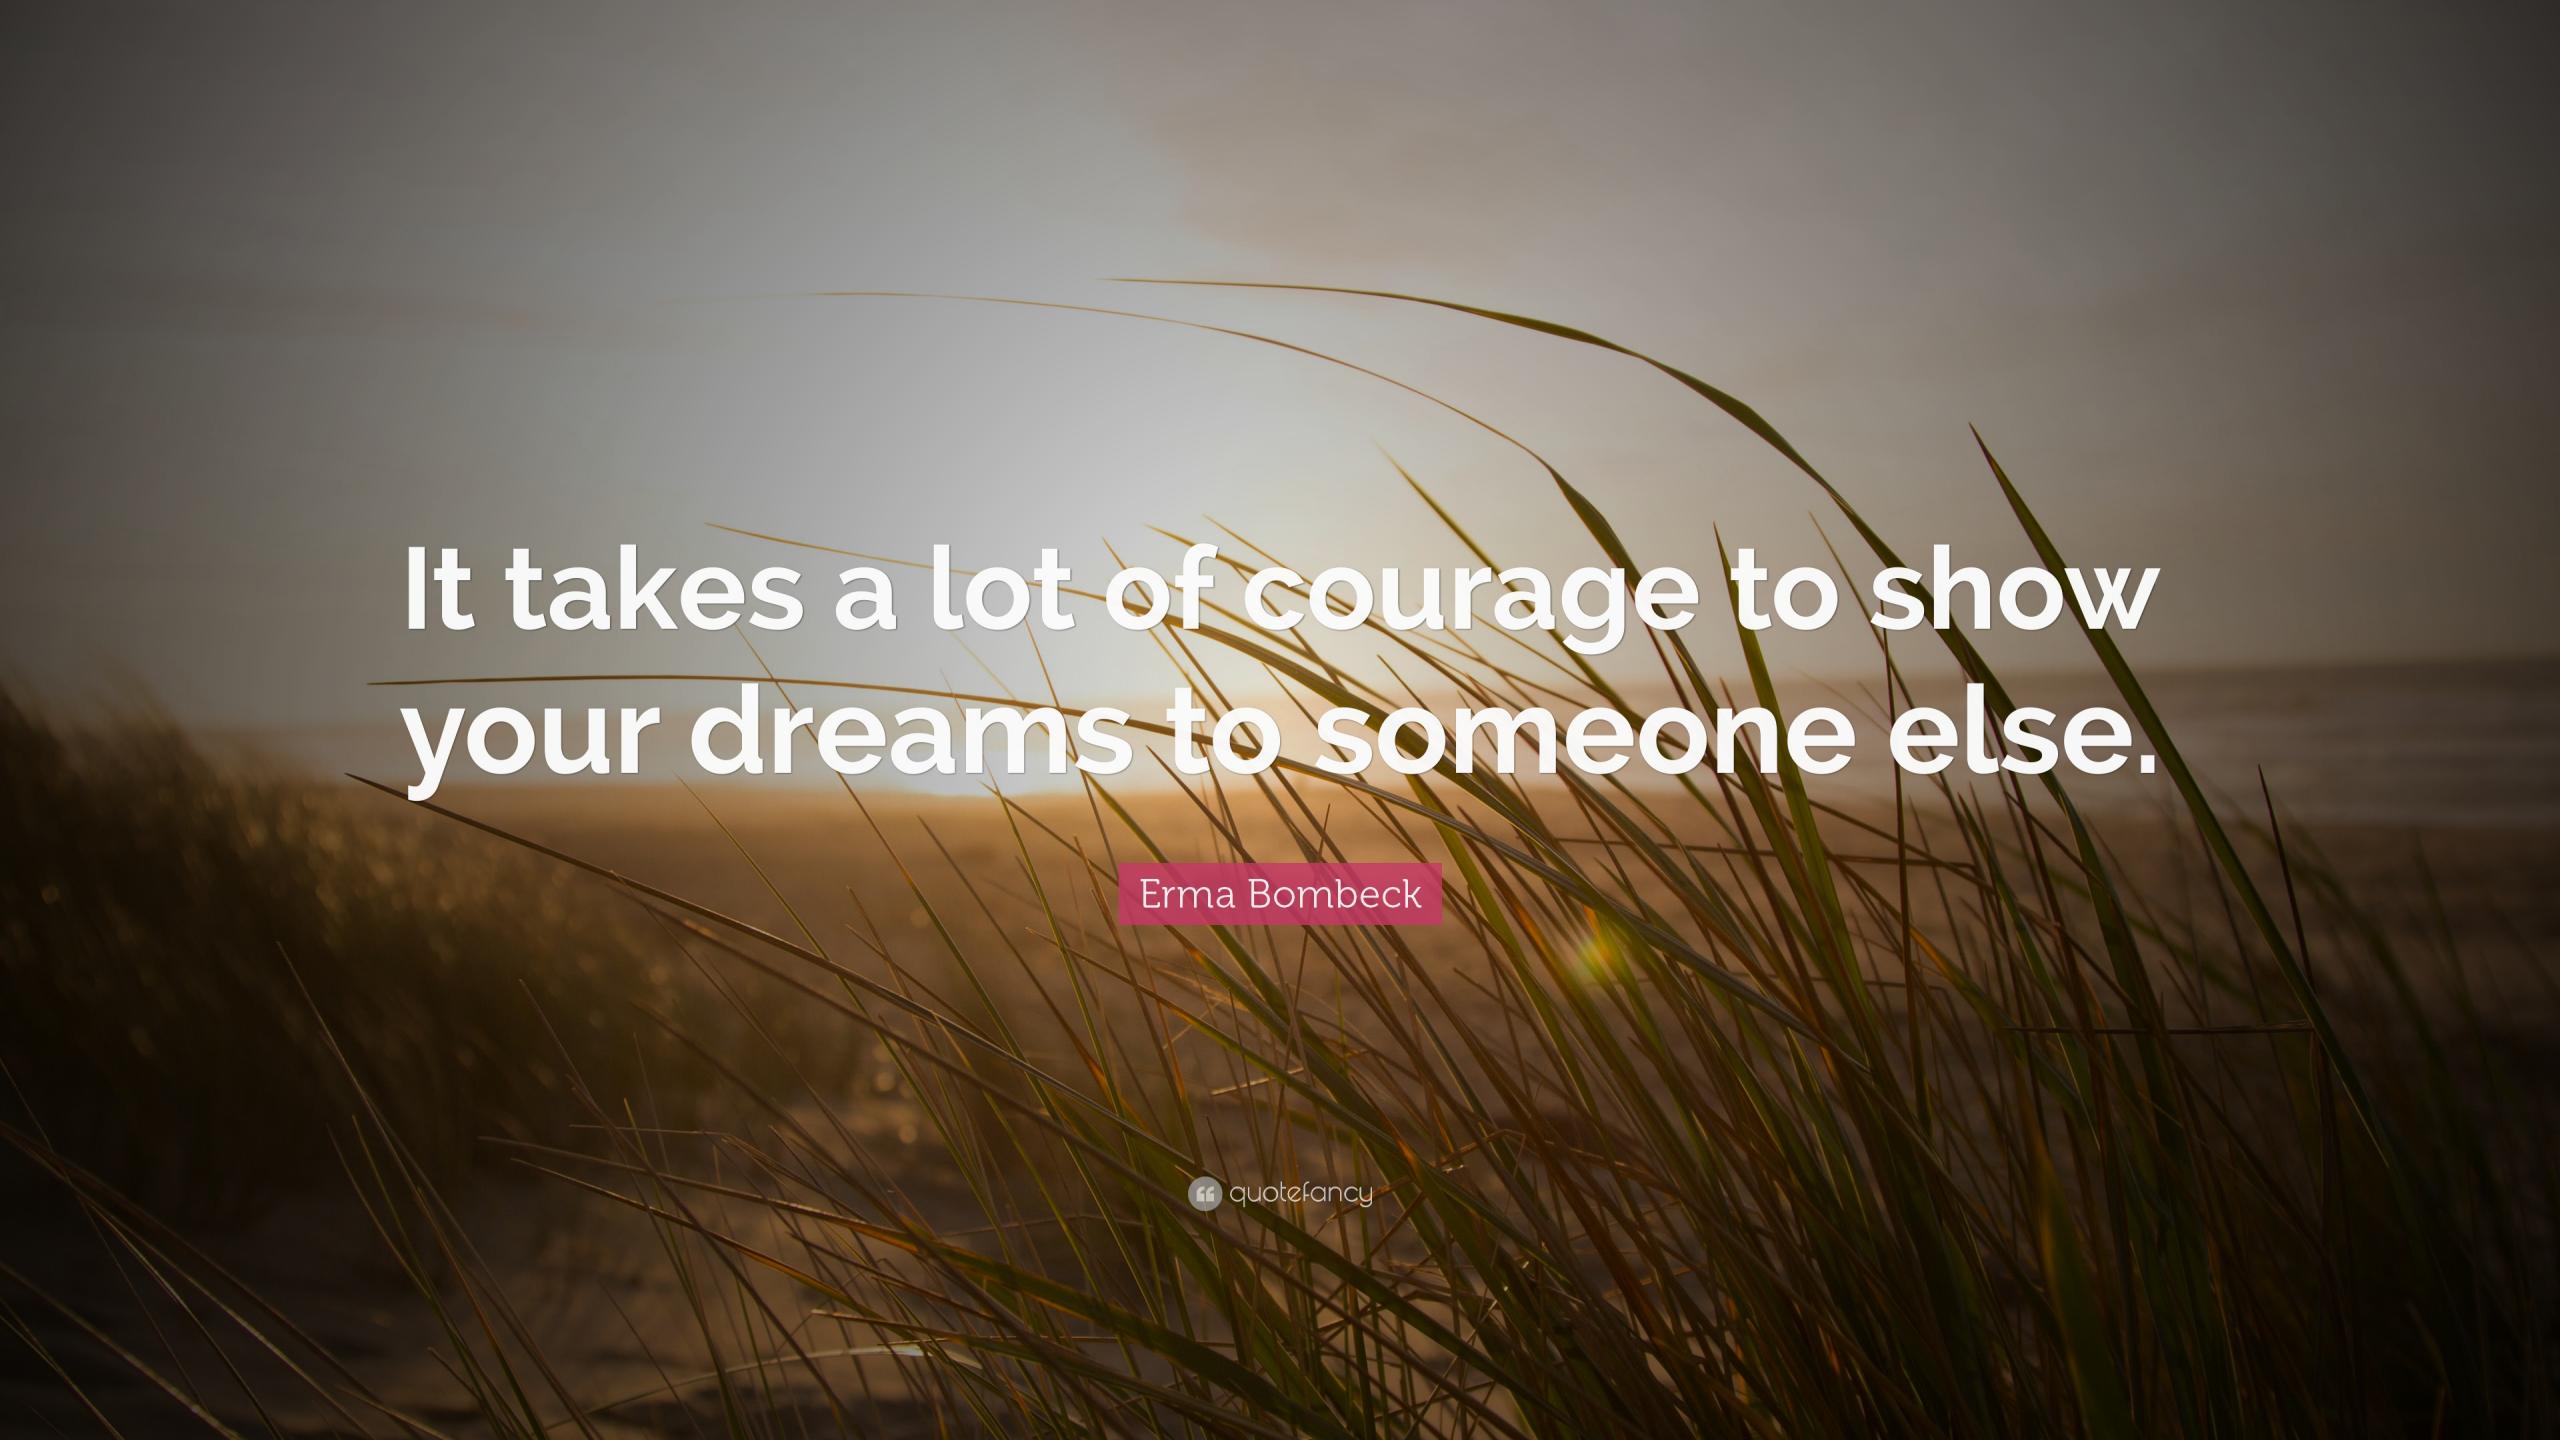
Unified Analytics Engine for Large-Scale Data Processing  
Speed, Easy to Use, Generality & Runs Everywhere



## Apache Druid

High Performance Real-Time Analytics Database  
Column-Oriented Storage with SQL Query [Apache Calcite]





**It takes a lot of courage to show  
your dreams to someone else.**

Erma Bombeck



# Use Case ~ Apache Spark & Data Warehouse [DW]



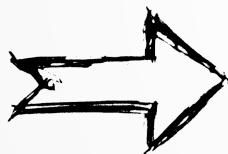
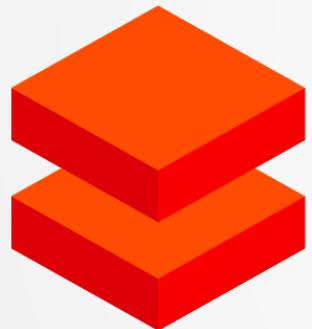
Distributed Cluster-Computing Framework  
Optimized for Memory Computation



MDW – Modern Data Warehouse  
MPP with Loosely Coupled Architecture



Report and Dashboards  
Push-Down Computation





**ONE WAY**  
SOLUTION